



## CPT-281 - Introduction to Data Structures with C++



### Module 15

### 2-3 Trees, 2-3-4 Trees, B-Trees

Dayu Wang

#### • 2-3 Trees

- These are **not binary** search trees.
- Because the nodes are not necessarily **binary**
  - ▶ They maintain all **leaves** at **same depth**
  - ▶ But number of children can vary
  - ▶ 2-3 tree: 2 or 3 children




2-node



3-node

### • Search in a 2-3 Tree

1. if r is NULL, return NULL (not in tree)
2. if r is a 2-node
  3. if target equals data1, return data1
  4. else if target < data1, search left subtree
  5. else search right subtree
6. else // r is a 3-node
  7. if target < data1, search left subtree
  8. else if target = data1, return data1
  9. else if target < data2, search middle subtree
  10. else if target = data2, return data 2
  11. else search right subtree




### • Inserting a key into a 2-3 tree

#### Inserting into a 2-Node Leaf

- **(Inserting 15)** We insert it directly, creating a new 3-node.

#### Inserting into a 2-Item Leaf with a 2-Node Parent


- **(Inserting 17)** The number will **virtually** be inserted into the 3 node at the bottom right
- Let's insert: 5, 10, and 20.



• Inserting a key into a 2-3 tree

Inserting into a 3-Node Leaf with a 3-Node Parent


- **(Inserting 13)** If we insert an item and all leaf nodes are full, one of the leaf nodes will need to be split.
- This would result in two new 2-nodes with values 10 and 13, and 11 would propagate up to virtually inserted in the 3-node at the root.
- Because the root is full, it would split into two new root node.



• Inserting a key into a 2-3 tree

Inserting into a 3-Node Leaf with a 3-Node Parent

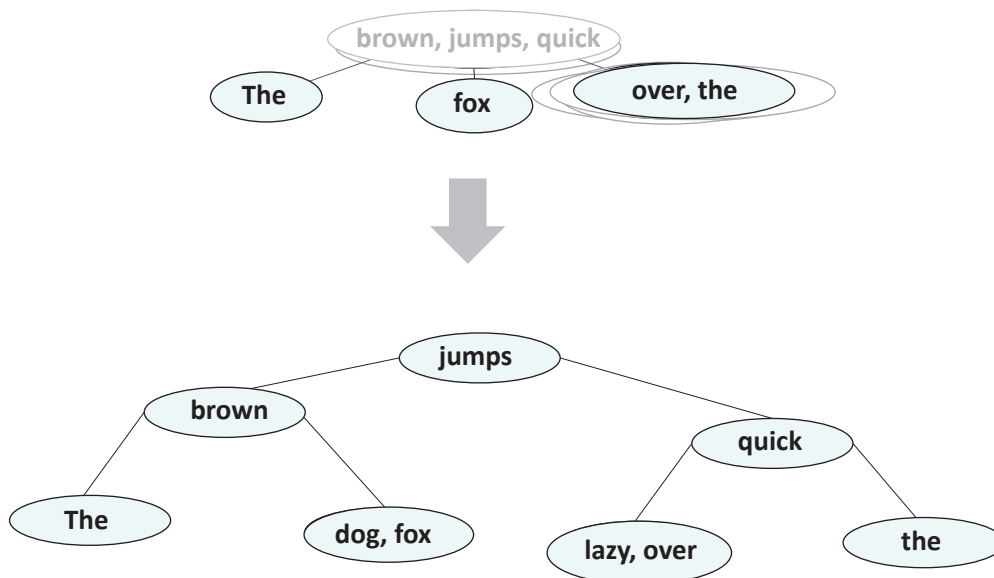
- If we insert an item and all leaf nodes are full, one of the leaf nodes will need to be split.
- This would result in two new 2-nodes with values 10 and 13, and 11 would propagate up to virtually inserted in the 3-node at the root.
- Because the root is full, it would split into two new root node.



• Inserting a key into a 2-3 tree

1. if  $r$  is NULL, return new 2-node with item as data
2. if item matches  $r \rightarrow \text{data1}$  or  $r \rightarrow \text{data2}$ , return **false**
3. if  $r$  is a leaf
4.     if  $r$  is a 2-node, expand to 3-node and return it
5.     split into two 2-nodes and pass them back up
6. **else**
7.     recursively insert into appropriate child tree
8.     if new parent passed back up
9.         if will be tree root, create and use new 2-node
10.        **else** recursively insert parent in  $r$
11. return **true**


• Building a 2-3 tree by inserting keys one-by-one




*"The quick  
brown fox  
jumps over  
the lazy  
dog"*

**• Exercise**

- Show how the 2-3 tree below as you insert “apple”, “cat”, and “hat” in that order.


**• Exercise**

- Adding “apple”



**• Exercise**


- Adding “cat” and hat

**• Exercise**

- Build a 2-3 tree from these numbers: 5, 18, 14, 16, 19, 25, 10

**• Exercise**

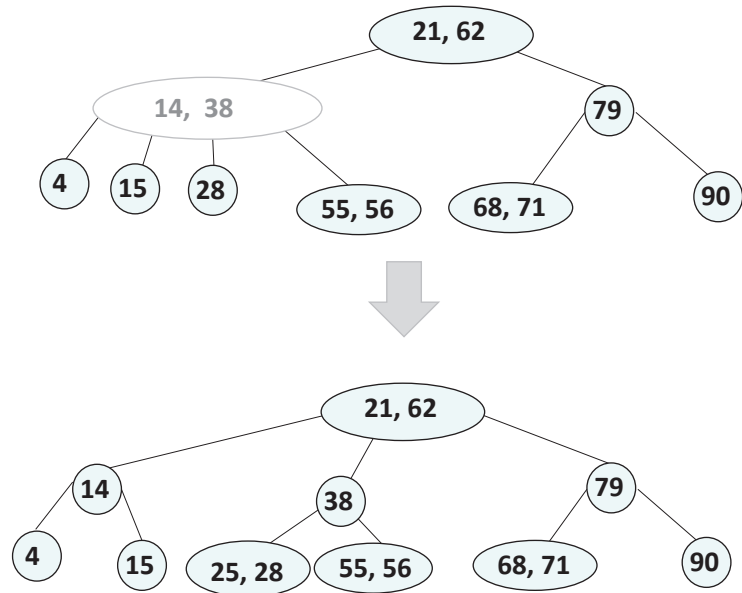
- Build a 2-3 tree from these numbers: 5, 18, 14, 16, 19, 25, 10

**• Performance of 2-3 trees**


- If height is  $h$ , number of nodes in range  $2^h - 1$  to  $3^h - 1$
- height in terms of # nodes  $n$  in range  $\log_2 n$  to  $\log_3 n$
- This is  $O(\log n)$ , since log base affects by constant factor
- So all operations are  $O(\log n)$

- Inserting keys into 2-3-4 trees

- **Inserting 25:** You search for where the item belongs according to the binary search tree algorithm. If a 4 node is encountered (a node with 3 values), you split it.
- **21** will propagate up to the parent.



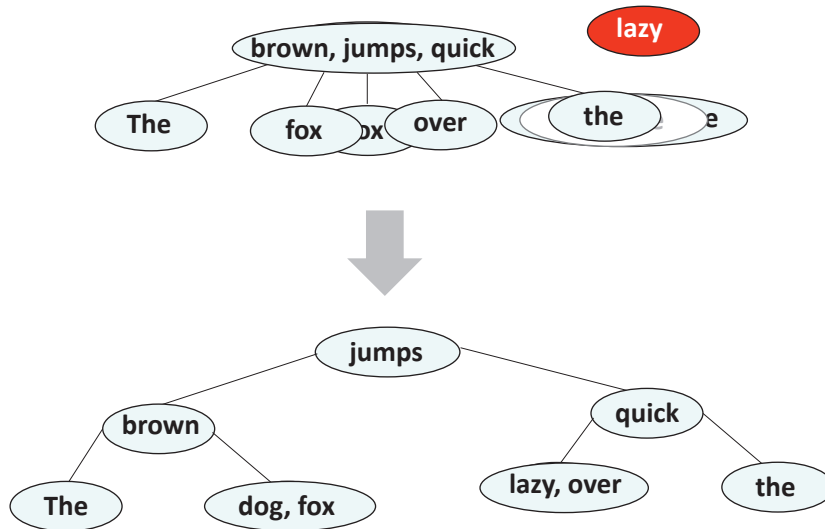
- Building 2-3-4 trees by inserting keys one-by-one



*"The quick  
brown fox  
jumps over  
the lazy  
dog"*




• Building 2-3-4 trees by inserting keys one-by-one







*"The quick  
brown fox  
jumps over  
the lazy  
dog"*

• Exercise

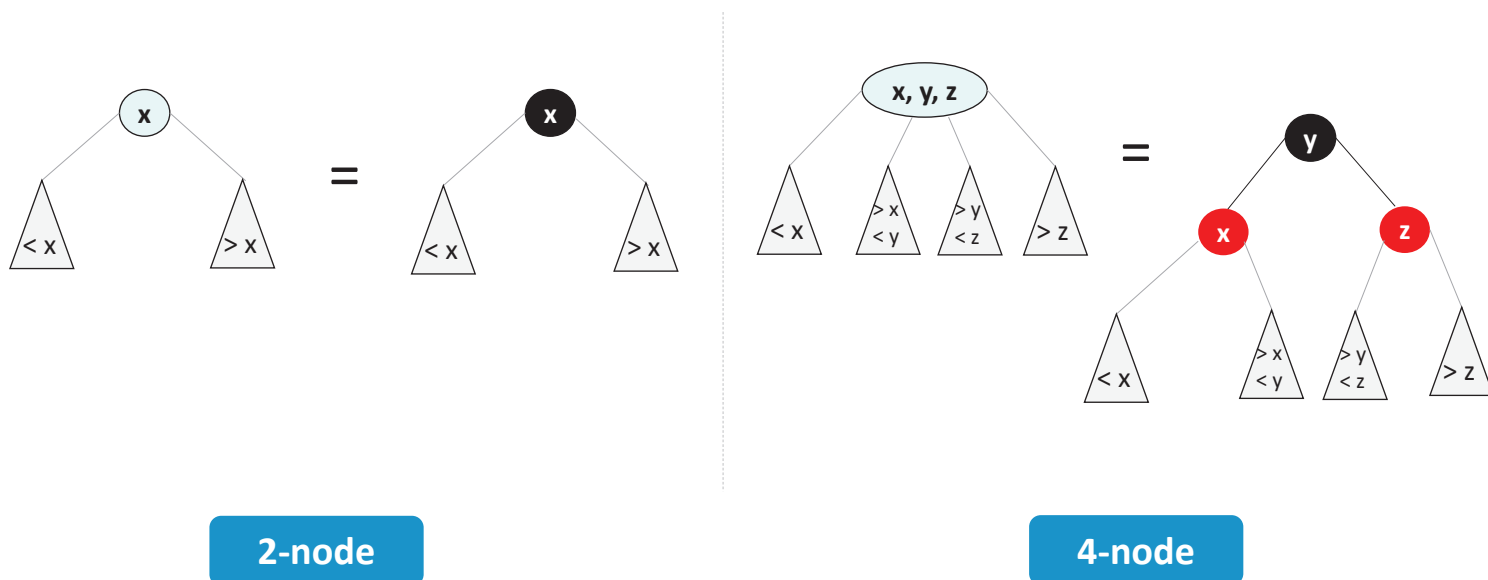
- Show the following tree after inserting each of the following values one at a time:  
1, 5, 9, and 13




• Exercise• Exercise

• Exercise• Exercise

• 2-3-4 Trees versus Red-Black Trees





• 2-3-4 Trees versus Red-Black Trees





### • 2-3-4 Trees versus Red-Black Trees

- We want to insert **3** into a 2-3-4 tree.



### • Exercise

- Convert the following 2-3-4 tree into a red-black tree



- External Study Resource

- 2-3 Trees (Insertion and Deletion)


<http://slady.net/java/bt/view.php>

- 2-3-4 Trees (Insertion and Deletion)

<http://www.cs.unm.edu/~rlpm/499/ttft.html>

- B-Trees

- A **B-tree** extends idea behind 2-3 and 2-3-4 trees:
- Allows a maximum of data items in each node
- **Order** of a B-tree is maximum number of children for a node
- B-trees developed for indexes to databases on disk



number of children: "5"

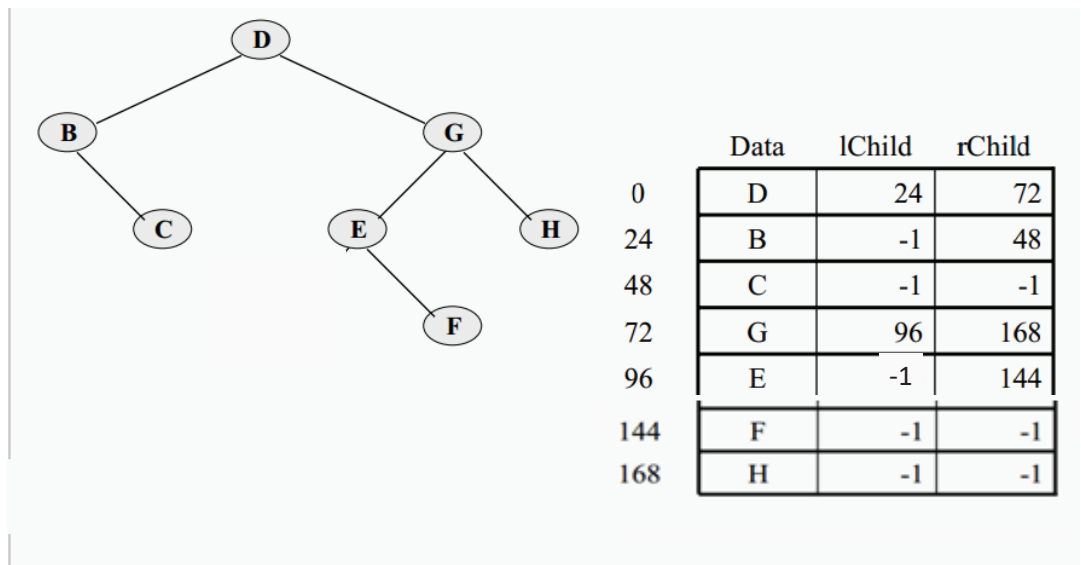
### • Motivation for B-Trees

- Index structures for large datasets cannot be stored in main memory
- Storing it on disk requires different approach to efficiency
- Assuming that a disk spins at 3600 RPM, one revolution occurs in 1/60 of a second, or 16.7ms
- Crudely speaking, one disk access takes about the same time as 200,000 instructions

Reference: <http://cecs.wright.edu/~tkprasad/teaching.html>

### • Storing tree data on a disk

Storing a tree on a disk



Reference: <http://courses.cs.vt.edu/cs2604/fall05/wmcquain/Notes/C12.B-Trees.pdf>


**• Motivation for B-Trees**

- Assume that we use an AVL tree to store about 20 million records
- We end up with a **very** deep binary tree with lots of different disk accesses;  $\log_2 20,000,000$  is about 24. Assuming each node is stored on a sector. The height of the tree is about 24. To find an item, we might need to navigate across 24 nodes to reach the target. so this takes about  $24 * (1/60)$  which is about 0.4 seconds.
- We know we can't improve on the  $\log n$  lower bound on search for a binary tree
- But, the solution is to use more branches and thus reduce the height of the tree!
- As branching increases, depth decreases

Reference: <http://cecs.wright.edu/~tkprasad/teaching.html>

**• Disk Storage**

- Disk storage is broken into blocks, and the time to access a block is significant compared to the time required to access data in internal memory.
- The nodes of a B-tree are sized to fit in a block, so each disk access to the index retrieves exactly one B-tree node.



[http://www.dba-oracle.com/t\\_history\\_disk.htm](http://www.dba-oracle.com/t_history_disk.htm)



### • Analysis of B-Trees

- The maximum number of items in a B-tree of order  $m$  and height  $h$ :

root  $m - 1$

level 1  $m(m - 1)$

level 2  $m^2(m - 1)$

...

level  $h$   $m^h(m - 1)$

- So, the total number of items is

$$(1 + m + m^2 + m^3 + \dots + m^h)(m - 1) =$$

$$[(m^{h+1} - 1) / (m - 1)] (m - 1) = \mathbf{m^{h+1} - 1}$$

- When  $m = 5$  and  $h = 2$  this gives  $5^3 - 1 = 124$

Reference: <http://cecs.wright.edu/~tkprasad/teaching.html>

### • Analysis of B-Trees

- When searching tables held on disc, the cost of each disc transfer is high but doesn't depend much on the amount of data transferred, especially if consecutive items are transferred:
  - If we use a B-tree of order 101, say, we can transfer each node in one disc read operation
  - A B-tree of order 101 and height 3 can hold  $101^4 - 1$  items (approximately 100 million) and any item can be accessed with 3 disc reads.

Reference: <http://cecs.wright.edu/~tkprasad/teaching.html>